

# CIS452 Lecture-40 Notes

Winter 2017

**Scribe:** *Tyler Paquet*

## Announcements

- Continued to look through handout on Block Allocation

## Chapter 11 (Continued)

### 1. Indexed Allocation: Limitations

- The maximum file size is the number of pointers that can fit into one index block
- To store large files
  - Link several index blocks together
  - Multi-level index
  - Combined: use both direct index and multilevel indices (Unix File Systems)

### 2. UFS File Format

- UFS (Unix File System) refers to a set of FS variants based on Unix design
- Boot block (usually a single sector): code to bootstrap the OS
- Super block (several disk sectors): metadata about the filesystem
- Inode block (several disk sectors) for storing index nodes (inodes)
  - One inode per file/directory
- Many data blocks (greater than 95 percent of the disk sectors)
  - Directory entry and File objects
  - Journals used during recovery

### 3. UFS Super Block

- Size of the entire filesystem
- Information on free data blocks
  - Number of free blocks
  - A list of free blocks
- Information on free inodes

- Number of free inodes
- A list of free inodes
- The inode for the root (/) directory (inode 128 = root)

#### 4. Linux/Unix File Allocation

- One inode per file / directory
- File block addresses are stored in inode and also in indirect blocks
- N (10 or 12) direct pointers and 3 indirect pointers
  - N direct pointers hold the addresses of the first N data blocks used by a file
  - The N+1st pointer holds the address of a single indirect block
  - The N+2nd pointer holds the address of a double indirect block
  - The N+3rd pointer holds the address of a triple indirect block

#### 5. Indirect Blocks

- Indirect blocks are data blocks used for storing pointers to data blocks
- Single indirect blocks
  - Contain pointers to data blocks
- Double indirect blocks
  - Contain pointers to single indirect blocks
- Triple indirect blocks
  - Contain pointers to double indirect blocks

#### 6. Free (Disk) Block Management

- Bit Vector/Bitmap
- Linked-List (Free List)
  - Simple
  - Grouping
  - Counting
- In the list can be implemented as a tree (for faster search)

#### 7. Bit Vector / Bitmap

- Each block is represented by 1-bit (0 = free, 1 = used)
- Number of bytes for bitmap = number of total disk blocks
  - A disk of 64 blocks = 64 bits (8 bytes) needed for bitmap
- Fast to locate free blocks: use bitwise operations
- Require extra space to maintain free blocks (too much wasted space)
  - A disk with capacity 512G (= 239) and block size of 512 bytes (= 29)
  - Number of blocks =  $239/29 = 230$  blocks

- Size of bitmaps  $230/23 = 227$  bytes (128M) of bitmap
- Percentage of wasted space  $= 227 / 239 = 2^{-12} = 1/4096 \approx 2.5$

#### 8. Linked-List of Free Blocks / Free List

- Improvements over Bitmap technique
  - Chain all the free blocks together
  - The superblock keeps the head and tail of this list
- Place the next pointer inside (at the end) of the freeblocks
- Searching for contiguous free blocks requires reading each block in the chain (too many disk I/O)

#### 9. Efficiency of Disk Block Operations

- Any bytes in RAM can be accessed within the same amount of time (ns)
- But, blocks on a disk may have different access time depending on their relative distance to the R/W head current position
- Updating just one byte on a disk block requires writing the entire block
- Design objectives of various FS are
  - higher storage capacity (how to support Petabytes or more)
  - fewer disk operations
  - highly consistent data (fewer recovery required)

#### 10. Critical On-Disk Data Structures

- A filesystem essentially holds (at least) three separate trees:
  - A tree of files and directories (user-owned)
  - A tree of free data blocks (maintained by OS)
  - A tree of free index blocks (maintained by OS)
- These trees originate from the superblock!
- Any updates to the FS must guarantee the three trees are in sync!

#### 11. Low-Level Filesystem Operations

- To write a file into a (Linux) filesystem
  - Write the file contents into the data blocks
  - Update the inode (file size, timestamps, block pointers, .)
  - Update the superblock (update the free-block list free-inode list)
- When the above sequence does not run to completion (i.e. by power failure), the filesystem records only a partial (inconsistent) data/metadata of your file

#### 12. When a new file is created

- Its data blocks are allocated [superblock update]

- One inode is allocated [superblock update and inode update]
- The directory entry for the file is update [data block update]
- The inode of the containing directory is updated [inode update]

### 13. Database Transactions (Intent) Logs

- To guarantee data integrity
- Multiple operations that update different parts of the DB are performed under one transaction: BEGIN TRANSACTION and COMMIT/ROLLBACK
  - In addition to changing the data, a transaction also logs the intended changes (insert, delete, update)
  - The log must be securely saved PRIOR TO the modification of the actual data
- At recovery time, the contents of the log are compared with the actual data and unfinished transactions can be recovered

### 14. Journaling / Log-Structured FS

- Every modification to the FS must be associated with a log entry
  - Log entry format: timestamp, action, blockid, oldvalue, newValue
  - Log entries are created for ALL types of modification (including inode and superblock)
- The log must be saved FIRST before the actual block updates take place
- Delete the log after the block updates are successful
- Recovery after system crash: use the existing log entries to restore FS